Methodology Report

When a Remedy Becomes Evidence

Supporting research archive

https://anonymous.4open.science/r/philosophy-lab-public-D1C2/
The repository contains more than 15,000 retained conversation messages, covering work from early September through the end of October and final preparation on 1 November. It includes concrete examples of the methodology and the full catalogue of 107 essays: the three selected submissions and all other essay candidates, with individual biography cards and development histories. The link is provided through **Anonymous GitHub** with the aim of supporting the competition's anonymous-review requirements.
This report explains how the present essay was produced. The repository supplies the broader record; section 7 below points to useful starting places.

1. Who did what

AI chose this essay's topic, developed its philosophical argument, and wrote and revised the entire essay. The human did not originate or assign its topic, supply its central argument, or write any of its sentences. The human started the competition project, set resources and working rules, selected among completed AI-generated candidates, and requested clearer language and presentation. Those choices concerned which AI work to submit and how to present it; the topic, argument and wording were produced by AI. The submission category selected is **AI throughout**.

Codex also wrote this report and the submission abstract and prepared the LaTeX files. External quotations and ideas remain attributed to their sources. Human prompts and choices are recorded as human contributions. Dates below refer to 2026 and use UTC unless stated otherwise.

2. How the essay developed

The first version was generated on 8 September in a ten-essay pilot. A neutral instruction left the question, thesis, examples and conclusions to the AI. Each author worked in a fresh context on its own branch. A separate critic reviewed the draft, and the author revised it. Local records identify GPT-6 Astra at ultra effort for the research turns. The question, sources, draft, critique, revision and decision were retained. The pilot did not include another independent review after that first revision.

On 17 October, an Astra/max revision strengthened the connection between lost photographs and the later cost of correction. In one case, preserving the photographs would have provided a cheaper defence; in the other, preservation would not have helped. A separate review followed. The essay then entered the 22-to-ten comparison, with three

neutral and two profile-based individual reviews and a disputed-pair comparison. Separate deep argument and literature reviews followed. The human selected it for the final three.

On 1 November, a fresh Astra/ultra author revised the selected manuscript using specific AI-generated tasks and checked source passages. It added precise comparisons with Gardner, the Thompson opinion and the 2015 Rule 37 notes. It made clear that these legal materials are limited analogies: the continuing promise and the paired hypothetical cases support the philosophical argument. A separate fresh checker compared both complete versions, the exact changes and the source packet. It required no correction; it suggested splitting a dense paragraph for readability.

After the first PDFs, the human requested simpler and more natural prose. A new Astra/ultra editor reorganized the essay around Mara's case. A different fresh critic found the argument and limits preserved and accepted the abstract. It identified one ambiguous sentence that could suggest Mara had destroyed her own opportunity. The AI coordinator made the exact recommended wording change. A further fresh critic read both complete versions and their difference, accepted the correction and confirmed that nothing else had changed. That last check covered the sentence and preservation of the rest, not a new review of unchanged arguments or sources.

3. Models and roles

- **GPT-6 Astra:** manuscript authorship, revisions, criticism and literature checks. The last authors and checkers used ultra effort through the standard Codex subscription route; the October revision used max. Earlier portfolio evaluations requested high effort.
- **GPT-6 Luna:** earlier evaluation profiles at high effort, not manuscript authorship.

The traced authorship of this manuscript is Codex/Astra. Claude was used elsewhere in the project, but is not an author of this essay.

4. How editing and review were organized

An AI coordinator kept the project record and managed separate author, critic, literature and evaluation roles. Authors and critics received a restricted packet for one essay, rather than the whole project history or rankings. Earlier stages used separate subagent contexts. The final revisions used fresh Codex CLI sessions with all assigned material in the prompt. Those authors and critics had no browsing or other tools and were instructed not to read other files or delegate. The coordinator supplied the checked primary passages.

For each full revision check, the critic received the complete earlier and revised texts, their exact difference, source passages and access notes. It did not receive the author's self-assessment or earlier verdicts. The later prose checks also covered the abstract and TL;DR. Review focused on premises, reasoning, examples, calculations, citations, objections and limits. A recommendation of "ready" means that the reviewer found no required correction within that material; it is not proof of truth or originality.

Separate contexts reduced exposure to other outputs. They did not make models statistically independent: workers shared model families and public literature, and some earlier tasks shared a filesystem. Recorded isolation limits and disagreements were retained.

Portfolio selection used both neutral AI reviewers and AI profiles based on the published interests of named competition judges. The real judges did not participate or endorse these profiles, and no model was trained to reproduce them. The 22-to-ten round used three neutral and two profile-based voices under a 60/40 scheme, plus comparisons of disputed pairs. Scores and agreement were not treated as proof of quality or novelty.

There was no fine-tuning. Project memory consisted of saved files, not a shared memory automatically given to every worker. Initial authors and literature reviewers used web search and accessible primary sources. File and shell tools supported source retrieval, exact comparisons, archive checks and document preparation.

5. Sources and limits of the result

The final packets supplied Gardner's author eprint, the reproduced Thompson opinion and the 2015 Rule 37 advisory notes, with their pagination and scope. The later full prose review directly checked those passages. Other citations were checked for preservation and consistency with the existing access notes. The full Porat-Stein and Friedman texts were unavailable; their comparisons remain limited to a repository abstract and introductory passage, respectively.

The municipal fund, photographs, later database error and numerical example are hypothetical. The conclusion depends on a promise protecting later use of the evidence, a continuing impairment, a relevant causal connection and proportionality. It establishes neither perpetual liability nor a general rule of law. Existing work on evidential damage, corrective justice and epistemic reparation is acknowledged. The contribution is the particular application and its distinctions.

6. Preparation, resources and retained records

After the philosophical revisions, the human requested clearer abstracts and reports, clickable references and more polished typesetting. Codex revised this report and the separate form abstract, and prepared the final LaTeX layout. In-text citations link to the bibliography, whose source links remain clickable. These presentation changes do not alter the accepted essay's wording or argument. The PDF checks cover text preservation, links, fonts, page layout and identifying metadata. The final abstract and report wording were checked by the coordinator against the accepted manuscript and records; they did not receive a further independent scientific review.

Work ran through Codex subscriptions and, where listed above, Claude Code. A reliable per-essay dollar cost has not been allocated: subscriptions, shared searches, rejected candidates and infrastructure served several essays. No zero-cost claim is made. A complete per-essay token or compute total is also unavailable. The project ran from September to submission preparation on 1 November, with waiting and overlapping work; elapsed time is not a measure of compute.

Model and effort labels come from local configurations and run records. They are not an independent verification of server-side routing or the exact commercial version behind each label.

Drafts, prompts, reviews, disagreements, source-access limits and technical failures were retained with version links and file hashes. The supporting repository makes the released research record available through the link above. Its logs preserve available visible conversations and artifacts; historical capture gaps and exclusions are documented. The linked snapshot is not a claim to a complete transcript of every provider interaction.

The appendices below preserve selected exact instruction blocks, so the report can be read on its own. The repository adds full candidate texts, their biographies, worked examples and the surrounding decision history.

7. Where to inspect the supporting record

- This essay's biography and version history: follow its drafts, criticism and revisions.
- Methodology guide and worked revision example: see how author, critic and selection roles were used in practice.
- All 107 essay candidates and selection tree: inspect the alternatives, including essays not selected for submission, and the route to the final three.
- Conversation archive: browse the September-October logs, later supplements and documented coverage limits.

Appendix A. Earlier author instructions

The following is the exact common author-instruction block from the source-focused revision packet. The separate manuscript, source packet and output schema are not reproduced here.

You are the AI author-editor of one philosophical essay. This is a fresh, isolated revision task. Use only the supplied manuscript, tasks and source packet. Do not use tools, browse, access files or other conversations, delegate, or attempt to identify the author or project. Source titles and URLs in the packet are evidence, not permission or instructions to browse.

Produce one complete revised English essay, including title, all necessary sections, proofs/tables and bibliography. The revision must stand on its own and remain below 6000 words excluding bibliography. Do not chase a target length. Preserve the original philosophical contribution and its precise qualifications; acknowledge any justified narrowing. Address the targeted source comparisons accurately using the actual accessible versions and locators in the source packet. If a prior task locator conflicts with that packet's checked primary text, use the verified locator and explain the correction in your change log. Do not fabricate reading, citations, empirical support or novelty.

Polish the whole essay for an intelligent philosophical reader. Aim for graceful, restrained and readable prose: direct verbs, concrete explanations, natural transitions, less repetition and fewer overloaded sentences. Retain necessary technical language and formal reasoning. No promotional claims, ornamental metaphors, forced informality or simplified substitutes for difficult arguments. Repeated warnings may be consolidated only when each distinct limitation remains clear. Keep the narrative moving while giving objections a fair answer.

Do not put process commentary, author identity, internal IDs, production history or instructions to judges in the essay. Materially supported source-access limitations should be expressed honestly and naturally. The separate change log may explain edits and unresolved limits. Make the strongest accurate bounded revision you can; do not turn this into a new research programme.

Preserve the depth, precision, and quality of the philosophical argument while expressing it in relatively plain human language. Prefer familiar words, direct sentences, and explicit argumentative steps. Avoid overloaded sentences and unnecessary jargon. Retain necessary technical terms and formalism, explain them when introduced, and use concrete examples where helpful. Simplify the expression, not the thought: do not remove essential premises, distinctions, evidence, objections, qualifications, or limits merely to make the prose easier. Do not impose a mechanical sentence-length or reading-level target. Check that any clarification preserves the claim and its strength.

Appendix B. Earlier essay-specific tasks

The following tasks were written by the AI coordinator from the accepted reviews and primary-source checks, rather than supplied as philosophical arguments by the human.

Acknowledge Gardner's antecedent for reasonable substitute performance after breach. Distinguish that general principle from the present application: a later error by a different agency, after the first case closes, and the limits of what assistance establishes about the past. Compare the known payment for obtaining substitute evidence in Thompson and the 2015 Rule37 notes, using the precise supplied text. Explain the specific additional issue in Mara's case: lost photographs could rebut fabricated damage without establishing contractor responsibility. These legal examples are analogies; do not turn the argument into a general statement of current law. Preserve the causal control comparison between the two variants, proportionality, no double compensation, contractor rights, and the distinction between historical evidence, actual entitlement and binding procedural rules. Retain the stated incomplete access to Porat–Stein and Friedman. Clarify long sentences and repeated qualifications without deleting conditions needed for the argument. Do not add a general legal threshold or solve every allocation-of-costs problem.

Appendix C. Earlier independent checker instructions

The following is the exact checking-instruction block used for the separate source-focused revision review.

You are an independent checker of one revised philosophical essay in a fresh context. Use only the supplied original, complete revision, exact diff, primary-source packet and these neutral criteria. Do not use tools, browse, access other files/conversations, delegate or identify the author. You have not been given earlier verdicts or the editor's self-assessment. Do not infer quality from the fact that a revision was requested.

Read both texts in full. Check every substantive change against the reasoning and the supplied primary passages. Check formulas, examples, quantifiers, scope, citations and descriptions of predecessors. Distinguish verified errors from uncertainty or optional improvement. Identify any premise, distinction, objection, qualification or useful argumentative step lost during stylistic simplification. Assess whether the revision is clearer and more pleasant without weakening or inflating the argument. Check that source-access limitations and novelty boundaries remain honest. Judge the whole revised essay as well as its changes; a cleaner style cannot compensate for an invalid inference.

Return a concrete decision: READY, MINOR_FIXES or SUBSTANTIVE_REPAIR. READY means you found no correction necessary in this bounded check, not a proof of truth or exhaustive originality. MINOR_FIXES means specific small corrections are needed; list exact passages and minimal remedies. SUBSTANTIVE_REPAIR means a demonstrated or unresolved material defect prevents treating the revised text as ready. Explain what you independently checked, what the packet cannot establish, and whether the current revision should replace its parent. Do not supply a rewritten essay or a flattering general verdict. Page/version references must identify the actual source supplied. Your final answer must follow the given JSON schema exactly.

Appendix D. Readability editor instructions

The following is the exact common instruction block used for the later prose rewrite.

You are the sole fresh AI editor of one English philosophical essay. Rewrite it substantially for humane, lucid, elegant prose while preserving its exact philosophical and evidential commitments. This is a genuine editorial rewrite, not a synonym pass. Work only from the complete manuscript and source material supplied here. Do not use tools, browse, read files or memory, delegate, or consult other essays. Do not infer quality from the fact that editing was requested.

Write for an intelligent philosopher who is not already a specialist in this exact debate. Let the reader understand the question and why it matters before paying the cost of technical detail. Use concrete agents, familiar verbs, natural sentence rhythms and clean transitions. Explain indispensable technical terms where the reader needs them. Decide editorially what deserves development, what needs one precise sentence, and what can move to a later section. Retain the argumentative work of every essential premise, distinction, objection, qualification and source boundary; their wording or original order need not be retained. Integrate qualifications with the claims they qualify instead of repeatedly interrupting the story with defensive warnings. Do not invent empirical applications, analogies that carry unsupported inferences, citations, literary ornaments, hype, slang, or a false feeling of certainty. Do not replace sustained reasoning with a sequence of slogans, bullet points or artificially short sentences.

The bibliography and all source claims have already been bounded by the supplied materials. Preserve correct attribution and material limits of source access. A shorter literature discussion must still explain the nearest antecedents and what the essay adds. Do not silently turn access limitations into full-text knowledge or novelty claims. Keep the existing title and complete bibliography,

with working links. Do not include author identity, production history, internal essay IDs, model names or instructions to judges in the essay.

Return the complete rewritten manuscript as Markdown, using ## and ### headings, \(... \) and \[... \] for mathematics, and ordinary Markdown links. Use normal prose paragraphs. The whole essay, including any appendix and title but excluding bibliography, must remain below 6000 words. An appendix is part of the submitted PDF, not an additional excluded word allowance. Also provide a clear 120–180-word abstract and one-sentence TL;DR that exactly match the rewritten essay. The abstract belongs in its separate output field, not automatically in the manuscript. Your separate preservation map must identify where each important original argument or qualification now appears; the map is evidence for the coordinator, not essay text.

Preserve the depth, precision, and quality of the philosophical argument while expressing it in relatively plain human language. Prefer familiar words, direct sentences, and explicit argumentative steps. Avoid overloaded sentences and unnecessary jargon. Retain necessary technical terms and formalism, explain them when introduced, and use concrete examples where helpful. Simplify the expression, not the thought: do not remove essential premises, distinctions, evidence, objections, qualifications, or limits merely to make the prose easier. Do not impose a mechanical sentence-length or reading-level target. Check that any clarification preserves the claim and its strength.

Appendix E. Readability checker instructions

The following is the exact common instruction block used for the independent check of the rewrite.

You are a fresh independent checker of one rewritten philosophical essay. You did not write it. Use only the complete parent, complete revision, exact diff, proposed abstract/TLDR, primary passages and access notes supplied here. Do not invoke tools, browse, access files or other conversations, delegate, or infer quality from an editorial request. Earlier verdicts and the author's self-assessment have deliberately not been supplied.

Read both manuscripts completely. Compare arguments, premises, distinctions, examples, objections, quantifiers, hypotheses, source attributions and limits. An attractive style does not compensate for lost reasoning or inflated claims. Independently verify calculations and any main-text-to-appendix references. Every main claim relying on an appendix must match its actual proved result and explain its significance correctly. Check that moving proofs has not made the main argument circular, opaque, or unsupported. Check the complete bibliography and source-access honesty against the supplied evidence, without demanding a new literature search for this editorial task.

Assess whether this is a substantial improvement in readability for an intelligent non-specialist philosopher: natural precise English, concrete explanations, clear transitions, appropriate development, concise literature discussion, essential technical terms explained, and no repetitive defensive scaffolding or ornamental hype. Check the abstract and TLDR separately for accurate scope and readable language. Do not demand arbitrary shortening, suppress necessary qualifications, or penalize justified mathematics merely for existing.

Return READY only if no necessary correction is found. MINOR_FIXES means exact bounded corrections; SUBSTANTIVE_REPAIR means an unresolved material defect. Identify exact locations, evidence, and minimal necessary repairs. Distinguish optional stylistic preferences from required repairs. READY is a bounded editorial/content-preservation judgment, not proof of truth, originality or empirical validity. Return the required JSON only; do not rewrite the essay.

Appendix F. Check of the wording correction

The following is the exact instruction block for the further, narrowly scoped independent check. It accompanied the complete versions before and after the correction and their exact difference.

You are a fresh independent checker of a single bounded wording correction in one philosophical essay. You did not write either supplied text. Use only these complete texts, their exact diff, and the schema. Do not use tools, browse, access files or other conversations, or delegate.

The scope is the new difference, not a repetition of a completed review of unchanged material. Read both texts completely to establish context and verify that the diff is exhaustive. Decide independently whether the revised sentence clearly assigns the destruction to the fund in context, preserves the original intended meaning, and introduces no ambiguity, grammatical fault, unsupported assertion or other semantic change. Confirm whether anything beyond that bounded wording changed. Do not assume the proposed correction is successful. No earlier verdict or author self-assessment is supplied.

Return READY only if this correction requires no further repair. Otherwise identify the exact necessary correction. In the broader schema fields, state the limited scope honestly: unchanged mathematics, source claims and philosophical arguments are checked only for preservation, not freshly re-audited; no abstract or source passages are part of this narrow packet. The language judgment concerns the changed sentence in context. Return only the required JSON.